


Introduction to Audio Content Analysis

Module 8.0: Classifiers

alexander lerch



introduction

overview

corresponding textbook section

Chapter 8: Musical Genre, Similarity, and Mood (pp. 155)

- **lecture content**

- training set and test set
- intuitive intro to machine learning
- classifier examples

- **learning objectives**

- describe the basic principles and challenges of data-driven machine learning approaches
- implement a kNN classifier in Matlab



introduction

overview

corresponding textbook section

Chapter 8: Musical Genre, Similarity, and Mood (pp. 155)

- **lecture content**

- training set and test set
- intuitive intro to machine learning
- classifier examples

- **learning objectives**

- describe the basic principles and challenges of data-driven machine learning approaches
- implement a kNN classifier in Matlab



classification

machine learning

- computer runs a 'generic' program
- adapts parameters to training data

⇒ **data driven** approach

data and its representation defines much of the outcome

- *validity*: is data representative sample and do features focus on important characteristics
- *reliability*: does data lead to accurate and consistent results
- *reproducibility*: will multiple runs result in similar results

classification

machine learning

- computer runs a 'generic' program
- adapts parameters to training data

⇒ **data driven** approach

data and its representation defines much of the outcome

- *validity*: is data representative sample and do features focus on important characteristics
- *reliability*: does data lead to accurate and consistent results
- *reproducibility*: will multiple runs result in similar results

classification

general steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 **(adjust** classifier settings, return to 4.)

classification

general steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 **(adjust** classifier settings, return to 4.)

classification

general steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 **(adjust** classifier settings, return to 4.)

classification

general steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 (adjust classifier settings, return to 4.)

classification

general steps

- 1 **define training set:** annotated results
- 2 **normalize** training set
- 3 **train** classifier
- 4 **evaluate** classifier with test (or validation) set
- 5 **(adjust** classifier settings, return to 4.)

classification

rules of thumb

● training set

- training set size vs. number of features
 - training set too small, feature number too large \Rightarrow *overfitting*
- training set **too noisy** \Rightarrow *underfitting*
- training set **not representative** \Rightarrow *bad classification performance*

● classifier

- classifier too complex \Rightarrow *overfitting*
- **poor classifier** \Rightarrow *bad classification performance*
 - \rightarrow different classifier

● features

- **poor features** \Rightarrow *bad classification performance*
 - \rightarrow new, better features
- **features not normalized** \Rightarrow *possibly bad classification performance*
 - feature distribution (range, mean, symmetry)

classification

rules of thumb

● training set

- training set size vs. number of features
 - training set too small, feature number too large \Rightarrow *overfitting*
- training set **too noisy** \Rightarrow *underfitting*
- training set **not representative** \Rightarrow *bad classification performance*

● classifier

- classifier too complex \Rightarrow *overfitting*
- **poor classifier** \Rightarrow *bad classification performance*
 - \rightarrow different classifier

● features

- *poor features \Rightarrow bad classification performance*
 - \rightarrow new, better features
- *features not normalized \Rightarrow possibly bad classification performance*
 - *feature distribution (range, mean, symmetry)*

classification

rules of thumb

● training set

- training set size vs. number of features
 - training set too small, feature number too large \Rightarrow *overfitting*
- training set **too noisy** \Rightarrow *underfitting*
- training set **not representative** \Rightarrow *bad classification performance*

● classifier

- classifier too complex \Rightarrow *overfitting*
- **poor classifier** \Rightarrow *bad classification performance*
 - \rightarrow different classifier

● features

- **poor features** \Rightarrow *bad classification performance*
 - \rightarrow new, better features
- features **not normalized** \Rightarrow possibly *bad classification performance*
 - feature distribution (range, mean, symmetry)

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - split training set into N parts (randomly, but preferably identical number per class)
 - select one part as test set
 - train the classifier with all observations from remaining $N - 1$ parts
 - compute the classification rate for the test set
 - repeat until all N parts have been tested
 - overall result: average classification rate

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification

evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classification


evaluation

- define **test set** for evaluation
 - test set *different* from training set
 - otherwise, same requirements
- example: **N -fold cross validation**
 - 1 split training set into N parts (randomly, but preferably identical number per class)
 - 2 select one part as test set
 - 3 train the classifier with all observations from remaining $N - 1$ parts
 - 4 compute the classification rate for the test set
 - 5 repeat until all N parts have been tested
 - 6 overall result: *average* classification rate

classifier examples

k-Nearest Neighbor (kNN)


- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors



classifier examples

k-Nearest Neighbor (kNN)


- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors



classifier examples

k-Nearest Neighbor (kNN)

- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors




$k = 3 \Rightarrow$ red majority

classifier examples

k-Nearest Neighbor (kNN)

- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors




$k = 5 \Rightarrow$ black majority

classifier examples

k-Nearest Neighbor (kNN)

- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors




$k = 7 \Rightarrow$ red majority

classifier examples

k-Nearest Neighbor (kNN)

- **training:** extract reference vectors from training set (keep class labels)
- **classification:** extract test vector and set class to majority of k nearest reference vectors
- **classifier data:** all training vectors



classifier examples



Gaussian Mixture Model (GMM)

- **training:** build model of each class distribution as superposition of Gaussian distributions
- **classification:** compute output of each Gaussian and select class with highest probability
- **classifier data:** per class per Gaussian: μ and covariance, mixture weight?

classifier examples

Gaussian Mixture Model (GMM)

- **training**: build model of each class distribution as superposition of Gaussian distributions
- **classification**: compute output of each Gaussian and select class with highest probability



● classifier data: per class per Gaussian: μ and covariance mixture weight?

classifier examples

Gaussian Mixture Model (GMM)


- **training:** build model of each class distribution as superposition of Gaussian distributions
- **classification:** compute output of each Gaussian and select class with highest probability
- **classifier data:** per class per Gaussian: μ and covariance, mixture weight?

classifier examples

Support Vector Machine (SVM)

- **training:**

- map features to high dimensional space




- find separating hyperplane through maximum distance of support vectors (data points)
- **classification:** apply feature transform and proceed with 'linear' classification
- **classifier data:** support vectors, kernel, kernel parameters

classifier examples

Support Vector Machine (SVM)

- **training:**

- map features to high dimensional space




- find separating hyperplane through maximum distance of support vectors (data points)
- **classification:** apply feature transform and proceed with 'linear' classification
- **classifier data:** support vectors, kernel, kernel parameters

classifier examples

Support Vector Machine (SVM)

- **training:**

- map features to high dimensional space



- find separating hyperplane through maximum distance of support vectors (data points)
- **classification:** apply feature transform and proceed with 'linear' classification
- **classifier data:** support vectors, kernel, kernel parameters

summary

lecture content

- **data-driven approach**

- general systems that learn behavior from data
- human interaction through
 - parametrization and procedures
 - data selection

- **training & test set**

- 1 must not overlap
- 2 must be representative

- **fine balance of inputs**

- 1 number of features
- 2 classifier complexity
- 3 amount and variability of data

